

Memory Allocators

Revision


- Virtual Memory
- Pointer Arithmetic
- Bitwise Operations

Contents

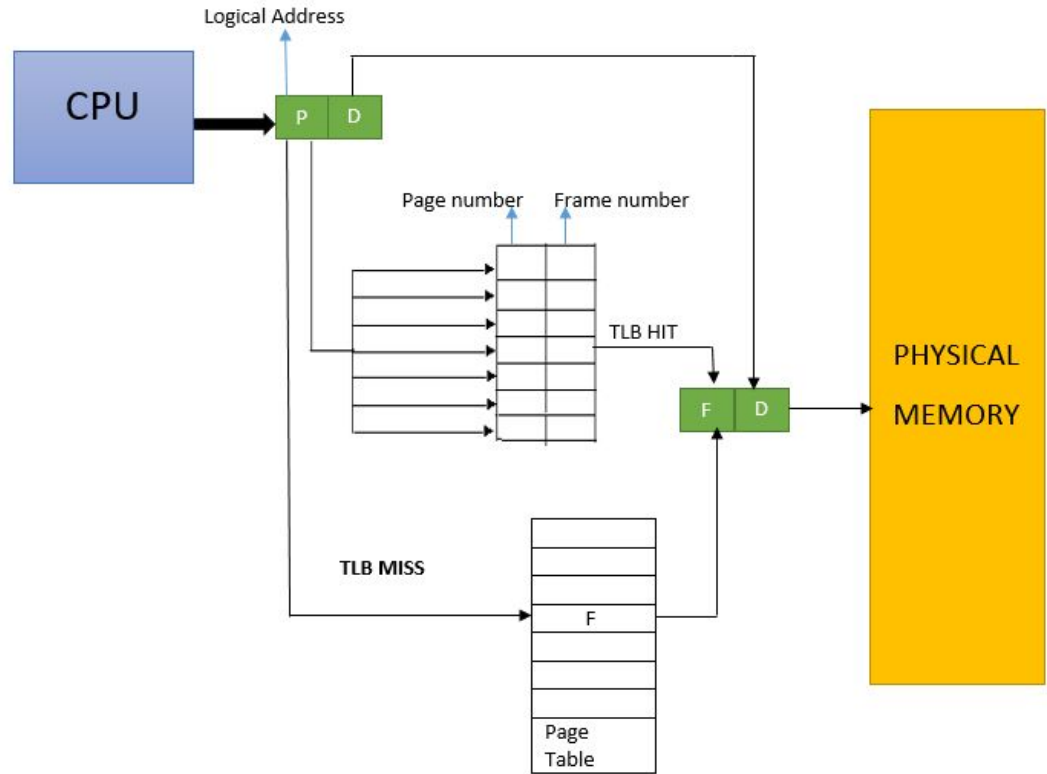
- Allocating Virtual Memory (Windows and Linux)
 - Memory Alignment
 - Types of Allocators
 - Bump (Arena) Allocator
 - Stack Allocator
 - Free list Allocator
 - Scratchpad/Temporary Memory
-

Revision

Virtual Memory



Virtual Memory



Allocating/Obtaining Virtual Memory

- Memory is a resource managed by the Operating System
- Use system calls to request OS for more memory
- The OS will allocate pages, creates mappings and returns back a pointer to memory

How much memory will the OS
return if we request 2 bytes?

Page Allocation

- OS will always allocated virtual memory in terms of pages.
- Consequence: Memory size allocated will always be a multiple of the Page size of the operating system
- Example, if page size of OS is 4kB, and we request a page having 2kB, then OS will return page of size 4kB regardless.

Memory Protections

- Provided by the OS.
- Three common types of protections:
 - Read Protection
 - Write Protection
 - Execute Protection

States of Memory

- Memory (in Windows) can be in 3 states: Reserved, Committed and Free
- Reserved
 - Only the address range is reserved and thus cannot be used by others
 - No pages are allocated
- Committed
 - Pages are allocated
- Free
 - The default state
 - Available for allocation

CRT Malloc:

A quick sneak under the hood

The “ugly” side of malloc


- Functions such as ``malloc`` and ``new`` are general-purpose memory allocators. Your code may be single-threaded, but the ``malloc`` function it is linked to can handle multithreaded paradigms just as well. It is this extra functionality that degrades the performance of these routines.
- Operating system has to _switch between user-space code and kernel code every time a request for memory/free is made.
- They form an entirely generic memory allocator. There is no restriction on the lifetime or the size of memory requested for allocation.

The mmap system call

- Used in Linux based system
- Maps files or devices into memory i.e, creates a memory mapping
- Enables memory mapped I/O

Memory Mapped I/O

- I/O in which devices/files are accessed through memory addresses (pointers)
- Common example is file I/O which utilizes your HDD or SSD
- Any changes made to the memory through the pointer will also be reflected into the file.



Relation with Virtual Memory

- Virtual pages are stored in disk
- *mmap* allows to creates mappings to files on disk
- Creating mappings to virtual pages
- Named as **anonymous mapping**

Using mmap for Memory Allocation

```
void *mmap( void *addr, size_t length, int prot, int flags, int fd,  
            off_t offset);
```

Returns a pointer to memory on success, MAP_FAILED on failure.

addr	Starting address of memory to allocate. If set to NULL, OS will return a suitable pointer instead.
length	Size of the memory in bytes we want to allocate.
prot	Protection flags
flags	Controls properties of the mapping
fd	Set to -1 for anonymous mapping
offset	Set to 0 for anonymous mapping

mmap: An Example

addr set to null

size of
4kB

Enable Read and
write to
allocated memory

```
size_t size = 4096;
void *mem = mmap( NULL, 4096, PROT_READ|PROT_WRITE,
                  MAP_ANON|MAP_PRIVATE, -1, 0 );

if ( mem == MAP_FAILED ){
    // .. error
    exit(1);
}
```

Create a
private and
anonymous map

fd and offset
set to -1 and 0

The munmap system call

- Undos a mapping created by mmap
- Basically, deallocates the pages allocated with mmap

```
int munmap( void *addr, size_t length );
```

- Undos the mapping of all pages that are in address range of **addr** to **addr+length**
- Returns 0 on success and -1 on error.

munmap Full Example

```
size_t size = 4096;
void *mem = mmap( NULL, size, PROT_READ|PROT_WRITE,
                  MAP_ANON|MAP_PRIVATE, -1, 0 );

if ( mem == MAP_FAILED ){
    // .. error
    exit(1);
}

if ( munmap( mem, size ) == -1 ){
    // error again
}
```

The VirtualAlloc function

- Windows counterpart of mmap
- “Reserve, commit or change the state of the region of memory within the virtual address space of a specified process.”

The VirtualAlloc system call

```
LPVOID VirtualAlloc(  
    [in, optional] LPVOID lpAddress,  
    [in]           SIZE_T dwSize,  
    [in]           DWORD  flAllocationType,  
    [in]           DWORD  flProtect  
);
```

- Needs header <Memoryapi.h>
- Returns an LPVOID type return value. LPVOID means Long Pointer to Void, i.e, void *
- Returns NULL when fails otherwise returns starting address of allocated memory

Side note on Windows Type Naming

LPVOID	Long Pointer to VOID type, i.e, a void *
SIZE_T	Equivalent to size_t
DWORD	A 32-bit unsigned integer
[in]	The parameter acts as an input to the function

The VirtualAlloc system call

```
LPVOID VirtualAlloc(  
    [in, optional] LPVOID lpAddress,  
    [in]           SIZE_T dwSize,  
    [in]           DWORD  flAllocationType,  
    [in]           DWORD  flProtect  
);
```

lpAddress	Similar to that of mmap
dwSize	Size of memory to be allocated
flAllocationType	Set to MEM_RESERVE MEM_COMMIT
flProtect	Protection Flags

VirtualAlloc: An Example

```
SIZE_T size = 4096;
void *mem = VirtualAlloc(NULL, size,
                        MEM_COMMIT|MEM_RESERVE, PAGE_READWRITE );

if ( mem == NULL ){
    //.... error
    exit(0);
}
// Success: Use mem
// .....
```


VirtualFree

- Decommits/Deallocates memory previously allocated by VirtualAlloc

```
BOOL VirtualFree(  
    [in] LPVOID lpAddress,  
    [in] SIZE_T dwSize,  
    [in] DWORD  dwFreeType  
);
```

lpAddress	Starting address of memory to be freed
dwSize	Size of memory to be freed
dwFreeType	MEM_DECOMMIT or MEM_RELEASE

States of Memory



MEM_DECOMMIT and MEM_RELEASE

- MEM_DECOMMIT sends pages back to MEM_RESERVED state
- MEM_RELEASE frees actual pages
 - dwSize **must** be 0 when MEM_RELEASE is used.

Virtual Free: An example

```
BOOL ret = VirtualFree(mem,0,MEM_RELEASE);  
if ( !ret ){  
    //... handle error  
}  
//.. No error
```

```
BOOL ret = VirtualFree(mem,4096,MEM_DECOMMIT);  
if ( !ret ){  
    //... handle error  
}  
  
ret = VirtualFree( mem, 0, MEM_RELEASE );  
//.. No error
```

So most of the times there is no
true “freedom”.

Revision

Memory Alignment

- A variable of a certain data type stored in memory follows an “alignment rule”
- The rule states which addresses are considered “valid” for a certain data type
- Some examples:
 - A 32-bit integer is always stored at an address that is a multiple of 4
 - A 64-bit integer is always stored at an address that is a multiple of 8
 - and so on..
- Dependent on the size of the data type.
- Access from an unaligned memory leads to slower code

Aligning to a Power of Two

```
uint64_t align_up( uint64_t size, uint64_t align ){  
    return ( size + ( align - 1 ) ) & ~(align-1);  
}
```

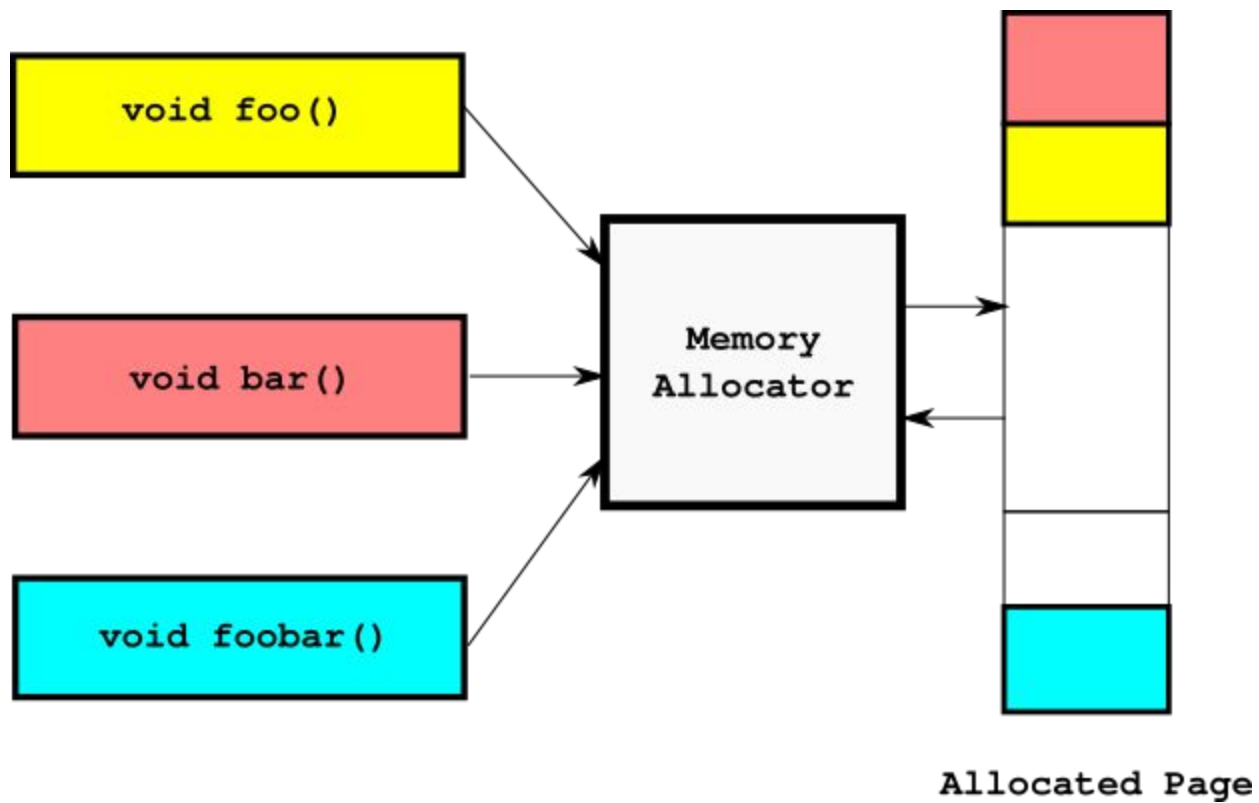
- The above functions “aligns” the value of size to the power of 2 given by align
- $\text{align_up}(23, 8) = 24$
- $\text{align_up}(27,8) = 32$
- $\text{align_up}(16,8) = 16$

Aligning Up to a power of Two: Breakdown

- What does $X \& \sim (Y-1)$ do when Y is a power of two?
 - Hint: Write $Y-1$ in binary
- What is the value of $(X + (\text{align}-1))$?
 - Check for different values of X , when X is a multiple of align and when it is not

Memory Allocators

- After allocating virtual memory from the OS, we are left with a single big chunk of memory
- We still need some way to distribute and manage this chunk of memory among different parts of our program
- Main functionalities:
 - An 'alloc' function to allow our program to request memory
 - A 'free' function to return memory back so that it can be reused.
 - Allows for addresses of arbitrary alignment



Memory Allocators: Types


- We discuss 4 different types of allocators
- In Increasing order of their complexity and flexibility:
 - Bump Allocator
 - Stack Allocator
 - Pooled style Free list Allocator
 - Generic Free List Allocator
- All allocators have their own way of performing memory allocation and deallocation

Bump Allocator

Bump Allocators

- Simplest kind of memory allocator
- Allows for memory allocations
- Allows limited freeing of memory
- It is **fast** and can be reused easily.

Bump Allocator: The Bigger Picture




- Different color indicate memory allocated by different parts of the program
- Gray region has address of 0x100, Red(maroon) has address of 0x200 and so on.
- The next allocation made will have the memory address 0x290.
- Alignment is not considered (yet).

Bump Allocator: Data Structure

- **Current** pointer, which points to base of unallocated memory
- Pointer to the memory allocated from the OS.
- Size of the memory allocated from the OS.

Bump Allocator: Data Structure

- In our previous example:
 - **Current** = 0x290
 - Pointer to memory allocated from OS = 0x100
 - Size of allocated memory = 1024 bytes



Bump Allocator: Allocation

- Consider allocation of X bytes of memory with a given alignment
- The steps involved include:
 - Aligning value of **current** pointer to the given alignment
 - Saving value of **current** pointer
 - Incrementing value of **current** pointer by X bytes
 - Performing bounds check
 - Returning the allocated memory (i.e, the save **current** pointer)
- Consider the allocation of 37 bytes initially, with an alignment of 8 bytes


Bump Allocator: Allocation

- Initial State:
 - **Current** = 0x100
 - Memory pointer = 0x100
 - Size = 1024 bytes
- Notice that **current** pointer is already aligned to 8 bytes




Bump Allocator: Allocation

- The **current** pointer is incremented by 37 bytes.
- Bounds checking the **current** pointer
 - Lies inside the allocated memory from the OS
 - Ensures that 37 bytes can be utilized
- The pointer value of 0x100 is returned




Bump Allocator: Second Allocation

- Now consider allocation of another 100 bytes with 8 byte alignment
- Align **current** pointer to 8 bytes
 - $\text{align}(0x125, 8) = 0x128$



Bump Allocator: Second Allocation

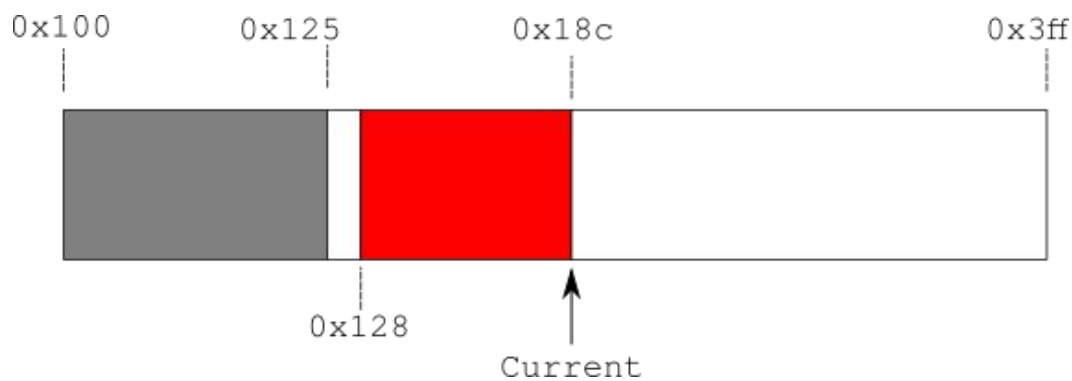
- Increment **current** pointer by 100 bytes
- Bounds checking the **current** pointer
 - Lies inside the allocated memory from the OS
- The pointer value of 0x128 is returned




Bump Allocator: Freeing Memory

- Limited when it comes to freeing memory
- Can only free memory from the top
- Size of memory to be freed needs to be passed
- Freeing is simple
 - Simply decrement the **current** pointer by the size to be freed
 - Bounds checking still needs to be done

Bump Allocator: Freeing Memory Example



After Freeing by 100 bytes:



Bump Allocator: Resetting (Important)

- Instead of freeing from the top, a more common operation is **resetting of the bump allocator**
- Simply sets the allocator to its initial state
- Easily Re-use the allocator for other purposes

Bump Allocator: Summary


- Advantages:
 - Simple and Fast
- Drawbacks
 - Cannot grow in size
 - Limited freeing
- Useful for temporary memory, i.e, applications where we need memory only for a relatively small number of operations. (*will come back to this later*)

Stack Allocator

Stack Allocators

- Similar to a Bump Allocator in their operation
- However, it keeps track of every allocation
- Freeing however, still can only be done from the top


Stack Allocators: The Bigger Picture



- Allocations are tracked using a linked list of pointers that goes backwards


Stack Allocator: Data Structure

- **Current** pointer, which points to base of unallocated memory
- Pointer to the memory allocated from the OS.
- Size of the memory allocated from the OS.
- **Top** pointer, which points to the beginning of the linked list structure




Stack Allocator: Allocation

1. Align the **current** point suitably so that we can store a pointer.




Stack Allocator: Allocation

2. Store the current **Top** value in that address




Stack Allocator: Allocation

3. Set the **Top** value to the **Current** pointer




Stack Allocator: Allocation

4. Point to the next memory address after. This is the address to be returned back. Any required alignment should be done in this step.



Stack Allocator: Allocation

5. Allocate the required number of bytes as done in bump allocator.




Stack Allocator: Freeing Memory

- Stack allocator provides a **Pop()** function that can be used to free memory at the top
- The **Pop()** does not need a size argument to be passed


Stack Allocator: Freeing Memory

1. Set **Current** = **Top**



Stack Allocator: Freeing Memory

2. Set $\text{Top} = \text{*Top}$



Stack Allocator: Summary

- Not very useful compared to bump allocator
- Added a lot complexity for addition of a very small functionality
- Total size of allocator is fixed.

Free List Allocator

Free List Allocator

- Uses a linked list to store memory chunks
- Memory chunks are of fixed size
- Allocations of only a fixed size can occur using this kind of free list
- Allows for general freeing of memory blocks
- We do NOT consider the case for an arbitrary aligned memory.

Free List Allocator (Pool-Based): Algorithm


1. Start with an empty page

0x200




Free List Allocator (Pool-Based): Algorithm

2. Create a linked list of memory chunks. Here, we have assumed a chunk size of 256 (0x100) bytes.




Free List Allocator (Pool-Based): Allocation

3. To allocate a chunk, remove the Top node from the linked list and return it.




Free List Allocator (Pool-Based): Allocation

Here's another example:



Free List Allocator (Pool-Based): Freeing Memory

4. To free memory, we simply insert it into the HEAD of the linked list and store the pointer to the next chunk.



Free List Allocator (Pool-Based): Getting New Memory

- What happens when free list runs out of chunks?
 - Simply allocate more Virtual Memory and perform the same steps as above
- Need to keep track of all the virtual memory allocated so far
 - Required to free all the created mappings

Question: Were the list headers
necessary? :)

When memory blocks are variable
sized!

Free List Allocator: Summary

- Allows for freeing of individual allocations
- Can grow in size
- Drawbacks
 - Allocation size is fixed (for pool-based)
 - For non-pool based free list, it is too generic.

Scratchpad/Temporary Memory

Scratchpad Memory

- Well technically it is a high-speed, low storage memory in hardware
- Used for **temporary storage to hold very small items** of data for rapid retrieval.
- Think cache but **without the cache replacement policy** and **more software control**.
- We'll use this concept to manage memory in our program.

Scratchpad Memory

- Well technically it is a high-speed, low storage memory in hardware
- Used for **temporary storage to hold very small items** of data for rapid retrieval.
- Think cache but **without the cache replacement policy** and **more software control**.
- Software implementation - remember what I said about resetting the scratchpad being important?

Temporary Allocation Region

Temporary String Allocation

Temporary Allocations inside Loop

Procedural Allocations

A thin horizontal line spans the width of the slide near the bottom. In the bottom-left corner, a diagonal line extends from the horizontal line towards the bottom-left corner of the slide.